

Python, Boto3, and AWS S3: Demystified

by Ralu Bolovan 20 Comments

[Tweet](#) [Share](#) [Email](#)

Table of Contents

- [Installation](#)
- [Client Versus Resource](#)
- [Common Operations](#)
 - [Creating a Bucket](#)
 - [Naming Your Files](#)
 - [Creating Bucket and Object Instances](#)
 - [Understanding Sub-resources](#)
 - [Uploading a File](#)
 - [Downloading a File](#)
 - [Copying an Object Between Buckets](#)
 - [Deleting an Object](#)
- [Advanced Configurations](#)
 - [ACL \(Access Control Lists\)](#)
 - [Encryption](#)
 - [Storage](#)
 - [Versioning](#)
- [Traversals](#)
 - [Bucket Traversal](#)
 - [Object Traversal](#)
- [Deleting Buckets and Objects](#)
 - [Deleting a Non-empty Bucket](#)
 - [Deleting Buckets](#)
- [Python Code or Infrastructure as Code \(IaC\)?](#)

- [Conclusion](#)
- [Further Reading](#)



Your Guided Tour Through the Python 3.9 Interpreter »

Watch Now This tutorial has a related video course created by the Real Python team. Watch it together with the written tutorial to deepen your understanding: [Python, Boto3, and AWS S3: Demystified](#)

Amazon Web Services (AWS) has become a leader in cloud computing. One of its core components is S3, the object storage service offered by AWS. With its impressive availability and durability, it has become the standard way to store videos, images, and data. You can combine S3 with other services to build infinitely scalable applications.

[Boto3](#) is the name of the Python SDK for AWS. It allows you to directly create, update, and delete AWS resources from your Python scripts.

If you've had some AWS exposure before, have your own AWS account, and want to take your skills to the next level by starting to use AWS services from within your Python code, then keep reading.

By the end of this tutorial, you'll:

- Be confident working with buckets and objects directly from your Python scripts
- Know how to avoid common pitfalls when using Boto3 and S3
- Understand how to set up your data from the start to avoid performance issues later
- Learn how to configure your objects to take advantage of S3's best features

Before exploring Boto3's characteristics, you will first see how to configure the SDK on your machine. This step will set you up for the rest of the tutorial.

Free Bonus: 5 Thoughts On Python Mastery, a free course for Python developers that shows you the roadmap and the mindset you'll need to take your Python skills to the next level.

Installation

To install Boto3 on your computer, go to your terminal and run the following:

Shell

```
$ pip install boto3
```

You've got the SDK. But, you won't be able to use it right now, because it doesn't know which AWS account it should connect to.

To make it run against your AWS account, you'll need to provide some valid credentials. If you already have an IAM user that has full permissions to S3, you can use those user's credentials (their access key and their secret access key) without needing to create a new user. Otherwise, the easiest way to do this is to create a new AWS user and then store the new credentials.

To create a new user, go to your AWS account, then go to *Services* and select *IAM*. Then choose *Users* and click on *Add user*.

Give the user a name (for example, *boto3user*). Enable *programmatic access*. This will ensure that this user will be able to work with any AWS supported SDK or make separate API calls:

Add user

Set user details

You can add multiple users at once with the same access type and permissions. [Learn more](#)

User name* [Add another user](#)

Select AWS access type

Select how these users will access AWS. Access keys and autogenerated passwords are provided in the last step. [Learn more](#)

Access type* **Programmatic access**
Enables an **access key ID** and **secret access key** for the AWS API, CLI, SDK, and other development tools.

AWS Management Console access
Enables a **password** that allows users to sign-in to the AWS Management Console.

* Required [Cancel](#) [Next: Permissions](#)

To keep things simple, choose the preconfigured *AmazonS3FullAccess* policy. With this policy, the new user will be able to have full control over S3. Click on *Next: Review*:

Add user

Set permissions

[Add user to group](#) [Copy permissions from existing user](#) [Attach existing policies directly](#) [Create policy](#)

[Filter policies](#) Showing 1 result

	Policy name	Type	Used as	Description
<input checked="" type="checkbox"/>	AmazonS3FullAccess	AWS managed	Permissions policy (9)	Provides full access to all buckets via the A...

[Cancel](#) [Previous](#) [Next: Review](#)

Select *Create user*:

Add user

Review

Review your choices. After you create the user, you can view and download the autogenerated password and access key.

User details

User name	boto3user
AWS access type	Programmatic access - with an access key
Permissions boundary	Permissions boundary is not set

Permissions summary

The following policies will be attached to the user shown above.

Type	Name
Managed policy	AmazonS3FullAccess



A new screen will show you the user's generated credentials. Click on the *Download .csv* button to make a copy of the credentials. You will need them to complete your setup.

Now that you have your new user, create a new file, `~/.aws/credentials`:

Shell

```
$ touch ~/.aws/credentials
```

Open the file and paste the structure below. Fill in the placeholders with the new user credentials you have downloaded:

Config File

```
[default]
aws_access_key_id = YOUR_ACCESS_KEY_ID
aws_secret_access_key = YOUR_SECRET_ACCESS_KEY
```

Save the file.

Now that you have set up these credentials, you have a `default` profile, which will be used by Boto3 to interact with your AWS account.

There is one more configuration to set up: the default region that Boto3 should interact with. You can check out the [complete table of the supported AWS regions](#). Choose the region that is closest to you. Copy your preferred region from the `Region` column. In my case, I am using `eu-west-1` (Ireland).

Create a new file, `~/.aws/config`:

Shell

```
$ touch ~/.aws/config
```

Add the following and replace the placeholder with the `region` you have copied:

Config File

```
[default]
region = YOUR_PREFERRED_REGION
```

Save your file.

You are now officially set up for the rest of the tutorial.

Next, you will see the different options Boto3 gives you to connect to S3 and other AWS services.

Client Versus Resource

At its core, all that Boto3 does is call AWS APIs on your behalf. For the majority of the AWS services, Boto3 offers two distinct ways of accessing these abstracted APIs:

- **Client:** low-level service access
- **Resource:** higher-level object-oriented service access

You can use either to interact with S3.

To connect to the low-level client interface, you must use Boto3's `client()`. You then pass in the name of the service you want to connect to, in this case, `s3`:

Python

```
import boto3
s3_client = boto3.client('s3')
```

To connect to the high-level interface, you'll follow a similar approach, but use `resource()`:

Python

```
import boto3
s3_resource = boto3.resource('s3')
```

You've successfully connected to both versions, but now you might be wondering, "Which one should I use?"

With clients, there is more programmatic work to be done. The majority of the client operations give you a [dictionary](#) response. To get the exact information that you need, you'll have to parse that dictionary yourself. With resource methods, the SDK does that work for you.

With the client, you might see some slight performance improvements. The disadvantage is that your code becomes less readable than it would be if you were using the resource. Resources offer a better abstraction, and your code will be easier to comprehend.

Understanding how the client and the resource are generated is also important when you're considering which one to choose:

- Boto3 generates the client from a JSON service definition file. The client's methods support every single type of interaction with the target AWS service.
- Resources, on the other hand, are generated from JSON resource definition files.

Boto3 generates the client and the resource from different definitions. As a result, you may find cases in which an operation supported by the client isn't offered by the resource. Here's the interesting part: you don't need to change your code to use the client everywhere. For that operation, you can access the client directly via the resource like so: `s3_resource.meta.client`.

One such client operation is `_generate_presigned_url()`, which enables you to give your users access to an object within your bucket for a set period of time, without requiring them to have AWS credentials.

Common Operations

Now that you know about the differences between clients and resources, let's start using them to build some new S3 components.

Creating a Bucket

To start off, you need an S3 [bucket](#). To create one programmatically, you must first choose a name for your bucket. Remember that this name must be unique throughout the whole AWS platform, as bucket names are DNS compliant.

If you try to create a bucket, but another user has already claimed your desired bucket name, your code will fail.

Instead of success, you will see the following error: `botocore.errorfactory.BucketAlreadyExists`.

You can increase your chance of success when creating your bucket by picking a random name. You can generate your own function that does that for you. In this implementation, you'll see how using the `uuid` module will help you achieve that. A UUID4's string representation is 36 characters long (including hyphens), and you can add a prefix to specify what each bucket is for.

Here's a way you can achieve that:

Python

```
import uuid
def create_bucket_name(bucket_prefix):
    # The generated bucket name must be between 3 and 63 chars long
    return ''.join([bucket_prefix, str(uuid.uuid4())])
```

You've got your bucket name, but now there's one more thing you need to be aware of: unless your region is in the United States, you'll need to define the region explicitly when you are creating the bucket. Otherwise you will get an `IllegalLocationConstraintException`.

To exemplify what this means when you're creating your S3 bucket in a non-US region, take a look at the code below:

Python

```
s3_resource.create_bucket(Bucket=YOUR_BUCKET_NAME,
                           CreateBucketConfiguration={
                               'LocationConstraint': 'eu-west-1'})
```

You need to provide both a bucket name and a bucket configuration where you must specify the region, which in my case is `eu-west-1`.

This isn't ideal. Imagine that you want to take your code and deploy it to the cloud. Your task will become increasingly more difficult because you've now hardcoded the region. You could refactor the region and transform it into an environment variable, but then you'd have one more thing to manage.

Luckily, there is a better way to get the region programmatically, by taking advantage of a `session` object. Boto3 will create the `session` from your credentials. You just need to take the region and pass it to `create_bucket()` as its `LocationConstraint` configuration. Here's how to do that:

Python

```
def create_bucket(bucket_prefix, s3_connection):
    session = boto3.session.Session()
    current_region = session.region_name
    bucket_name = create_bucket_name(bucket_prefix)
    bucket_response = s3_connection.create_bucket(
        Bucket=bucket_name,
        CreateBucketConfiguration={
            'LocationConstraint': current_region})
    print(bucket_name, current_region)
    return bucket_name, bucket_response
```

The nice part is that this code works no matter where you want to deploy it: locally/EC2/Lambda. Moreover, you don't need to hardcode your region.

As both the client and the resource create buckets in the same way, you can pass either one as the `s3_connection` parameter.

You'll now create two buckets. First create one using the client, which gives you back the `bucket_response` as a dictionary:

Python

>>>

```
>>> first_bucket_name, first_response = create_bucket(
...     bucket_prefix='firstpythonbucket',
...     s3_connection=s3_resource.meta.client)
firstpythonbucket7250e773-c4b1-422a-b51f-c45a52af9304 eu-west-1

>>> first_response
{'ResponseMetadata': {'RequestId': 'E1DCFE71EDE7C1EC', 'HostId': 'r3AP32NQk9dvbHSEPIbyYADT769VQEN/+xT2BPM6HCn'}}
```

Then create a second bucket using the resource, which gives you back a Bucket instance as the `bucket_response`:

Python

>>>

```
>>> second_bucket_name, second_response = create_bucket(
...     bucket_prefix='secondpythonbucket', s3_connection=s3_resource)
secondpythonbucket2d5d99c5-ab96-4c30-b7f7-443a95f72644 eu-west-1

>>> second_response
s3.Bucket(name='secondpythonbucket2d5d99c5-ab96-4c30-b7f7-443a95f72644')
```

You've got your buckets. Next, you'll want to start adding some files to them.

Naming Your Files

You can name your objects by using standard file naming conventions. You can use any valid name. In this article, you'll look at a more specific case that helps you understand how S3 works under the hood.

If you're planning on hosting a large number of files in your S3 bucket, there's something you should keep in mind. If all your file names have a deterministic prefix that gets repeated for every file, such as a timestamp format like "YYYY-MM-DDThh:mm:ss", then you will soon find that you're running into [performance issues](#) when you're trying to interact with your bucket.

This will happen because S3 takes the prefix of the file and maps it onto a partition. The more files you add, the more will be assigned to the same partition, and that partition will be very heavy and less responsive.

What can you do to keep that from happening?

The easiest solution is to randomize the file name. You can imagine many different implementations, but in this case, you'll use the trusted `uuid` module to help with that. To make the file names easier to read for this tutorial, you'll be taking the first six characters of the generated number's hex representation and concatenate it with your base file name.

The helper function below allows you to pass in the number of bytes you want the file to have, the file name, and a sample content for the file to be repeated to make up the desired file size:

Python

```
def create_temp_file(size, file_name, file_content):
    random_file_name = ''.join([str(uuid.uuid4().hex[:6]), file_name])
    with open(random_file_name, 'w') as f:
        f.write(str(file_content) * size)
    return random_file_name
```

Create your first file, which you'll be using shortly:

Python

```
first_file_name = create_temp_file(300, 'firstfile.txt', 'f')
```

By adding randomness to your file names, you can efficiently distribute your data within your S3 bucket.

Creating Bucket and Object Instances

The next step after creating your file is to see how to integrate it into your S3 workflow.

This is where the resource's classes play an important role, as these abstractions make it easy to work with S3.

By using the resource, you have access to the high-level classes (`Bucket` and `Object`). This is how you can create one of each:

Python

```
first_bucket = s3_resource.Bucket(name=first_bucket_name)
first_object = s3_resource.Object(
    bucket_name=first_bucket_name, key=first_file_name)
```

The reason you have not seen any errors with creating the `first_object` variable is that Boto3 doesn't make calls to AWS to create the reference. The `bucket_name` and the `key` are called identifiers, and they are the necessary parameters to create an `Object`. Any other attribute of an `Object`, such as its size, is lazily loaded. This means that for Boto3 to get the requested attributes, it has to make calls to AWS.

Understanding Sub-resources

`Bucket` and `Object` are sub-resources of one another. Sub-resources are methods that create a new instance of a child resource. The parent's identifiers get passed to the child resource.

If you have a `Bucket` variable, you can create an `Object` directly:

Python

```
first_object_again = first_bucket.Object(first_file_name)
```

Or if you have an `Object` variable, then you can get the `Bucket`:

Python

```
first_bucket_again = first_object.Bucket()
```

Great, you now understand how to generate a `Bucket` and an `Object`. Next, you'll get to upload your newly generated file to S3 using these constructs.

Uploading a File

There are three ways you can upload a file:

- From an `Object` instance
- From a `Bucket` instance
- From the `client`

In each case, you have to provide the `Filename`, which is the path of the file you want to upload. You'll now explore the three alternatives. Feel free to pick whichever you like most to upload the `first_file_name` to S3.

Object Instance Version

You can upload using an `Object` instance:

Python

```
s3_resource.Object(first_bucket_name, first_file_name).upload_file(
    Filename=first_file_name)
```

Or you can use the `first_object` instance:

Python

```
first_object.upload_file(first_file_name)
```

Bucket Instance Version

Here's how you can upload using a Bucket instance:

Python

```
s3_resource.Bucket(first_bucket_name).upload_file(  
    Filename=first_file_name, Key=first_file_name)
```

Client Version

You can also upload using the `client`:

Python

```
s3_resource.meta.client.upload_file(  
    Filename=first_file_name, Bucket=first_bucket_name,  
    Key=first_file_name)
```

You have successfully uploaded your file to S3 using one of the three available methods. In the upcoming sections, you'll mainly work with the `Object` class, as the operations are very similar between the `client` and the `Bucket` versions.

Downloading a File

To download a file from S3 locally, you'll follow similar steps as you did when uploading. But in this case, the `Filename` parameter will map to your desired local path. This time, it will download the file to the `tmp` directory:

Python

```
s3_resource.Object(first_bucket_name, first_file_name).download_file(  
    f'/tmp/{first_file_name}') # Python 3.6+
```

You've successfully downloaded your file from S3. Next, you'll see how to copy the same file between your S3 buckets using a single API call.

Copying an Object Between Buckets

If you need to copy files from one bucket to another, Boto3 offers you that possibility. In this example, you'll copy the file from the first bucket to the second, using `.copy()`:

Python

```
def copy_to_bucket(bucket_from_name, bucket_to_name, file_name):  
    copy_source = {  
        'Bucket': bucket_from_name,  
        'Key': file_name  
    }  
    s3_resource.Object(bucket_to_name, file_name).copy(copy_source)  
  
copy_to_bucket(first_bucket_name, second_bucket_name, first_file_name)
```

Note: If you're aiming to replicate your S3 objects to a bucket in a different region, have a look at [Cross Region Replication](#).

Deleting an Object

Let's delete the new file from the second bucket by calling `.delete()` on the equivalent `object` instance:

Python

```
s3_resource.Object(second_bucket_name, first_file_name).delete()
```

You've now seen how to use S3's core operations. You're ready to take your knowledge to the next level with more complex characteristics in the upcoming sections.

Advanced Configurations

In this section, you're going to explore more elaborate S3 features. You'll see examples of how to use them and the benefits they can bring to your applications.

ACL (Access Control Lists)

Access Control Lists (ACLs) help you manage access to your buckets and the objects within them. They are considered the legacy way of administrating permissions to S3. Why should you know about them? If you have to manage access to individual objects, then you would use an Object ACL.

By default, when you upload an object to S3, that object is private. If you want to make this object available to someone else, you can set the object's ACL to be public at creation time. Here's how you upload a new file to the bucket and make it accessible to everyone:

Python

```
second_file_name = create_temp_file(400, 'secondfile.txt', 's')
second_object = s3_resource.Object(first_bucket.name, second_file_name)
second_object.upload_file(second_file_name, ExtraArgs={
    'ACL': 'public-read'})
```

You can get the `ObjectAcl` instance from the `Object`, as it is one of its sub-resource classes:

Python

```
second_object_acl = second_object.Acl()
```

To see who has access to your object, use the `grants` attribute:

Python

>>>

```
>>> second_object_acl.grants
[{'Grantee': {'DisplayName': 'name', 'ID': '24aafdc2053d49629733ff0141fc9fede3bf77c7669e4fa2a4a861dd5678f4b5'}}
```

You can make your object private again, without needing to re-upload it:

Python

>>>

```
>>> response = second_object_acl.put(ACL='private')
>>> second_object_acl.grants
[{'Grantee': {'DisplayName': 'name', 'ID': '24aafdc2053d49629733ff0141fc9fede3bf77c7669e4fa2a4a861dd5678f4b5'}}
```

You have seen how you can use ACLs to manage access to individual objects. Next, you'll see how you can add an extra layer of security to your objects by using encryption.

Note: If you're looking to split your data into multiple categories, have a look at [tags](#). You can grant access to the objects based on their tags.

Encryption

With S3, you can protect your data using encryption. You'll explore server-side encryption using the AES-256 algorithm where AWS manages both the encryption and the keys.

Create a new file and upload it using `ServerSideEncryption`:

Python

```
third_file_name = create_temp_file(300, 'thirdfile.txt', 't')
third_object = s3_resource.Object(first_bucket_name, third_file_name)
third_object.upload_file(third_file_name, ExtraArgs={
    'ServerSideEncryption': 'AES256'})
```

You can check the algorithm that was used to encrypt the file, in this case AES256:

Python

>>>

```
>>> third_object.server_side_encryption
'AES256'
```

You now understand how to add an extra layer of protection to your objects using the AES-256 server-side encryption algorithm offered by AWS.

Storage

Every object that you add to your S3 bucket is associated with a [storage class](#). All the available storage classes offer high durability. You choose how you want to store your objects based on your application's performance access requirements.

At present, you can use the following storage classes with S3:

- **STANDARD**: default for frequently accessed data
- **STANDARD_IA**: for infrequently used data that needs to be retrieved rapidly when requested
- **ONEZONE_IA**: for the same use case as STANDARD_IA, but stores the data in one Availability Zone instead of three
- **REDUCED_REDUNDANCY**: for frequently used noncritical data that is easily reproducible

If you want to change the storage class of an existing object, you need to recreate the object.

For example, reupload the `third_object` and set its storage class to `Standard_IA`:

Python

```
third_object.upload_file(third_file_name, ExtraArgs={
    'ServerSideEncryption': 'AES256',
    'StorageClass': 'STANDARD_IA'})
```

Note: If you make changes to your object, you might find that your local instance doesn't show them. What you need to do at that point is call `.reload()` to fetch the newest version of your object.

Reload the object, and you can see its new storage class:

Python

>>>

```
>>> third_object.reload()
>>> third_object.storage_class
'STANDARD_IA'
```

Note: Use [LifeCycle Configurations](#) to transition objects through the different classes as you find the need for them. They will automatically transition these objects for you.

Versioning

You should use versioning to keep a complete record of your objects over time. It also acts as a protection mechanism against accidental deletion of your objects. When you request a versioned object, Boto3 will retrieve the latest version.

When you add a new version of an object, the storage that object takes in total is the sum of the size of its versions. So if you're storing an object of 1 GB, and you create 10 versions, then you have to pay for 10GB of storage.

Enable versioning for the first bucket. To do this, you need to use the `BucketVersioning` class:

Python

```
def enable_bucket_versioning(bucket_name):
    bkt_versioning = s3_resource.BucketVersioning(bucket_name)
    bkt_versioning.enable()
    print(bkt_versioning.status)
```

Python

>>>

```
>>> enable_bucket_versioning(first_bucket_name)
Enabled
```

Then create two new versions for the first file `Object`, one with the contents of the original file and one with the contents of the third file:

Python

```
s3_resource.Object(first_bucket_name, first_file_name).upload_file(
    first_file_name)
s3_resource.Object(first_bucket_name, first_file_name).upload_file(
    third_file_name)
```

Now reupload the second file, which will create a new version:

Python

```
s3_resource.Object(first_bucket_name, second_file_name).upload_file(
    second_file_name)
```

You can retrieve the latest available version of your objects like so:

Python

>>>

```
>>> s3_resource.Object(first_bucket_name, first_file_name).version_id
'eQgH6IC1VGcn7eXZ_.ayqm6Ndjh0Adv'
```

In this section, you've seen how to work with some of the most important S3 attributes and add them to your objects. Next, you'll see how to easily traverse your buckets and objects.

Traversals

If you need to retrieve information from or apply an operation to all your S3 resources, Boto3 gives you several ways to iteratively traverse your buckets and your objects. You'll start by traversing all your created buckets.

Bucket Traversal

To traverse all the buckets in your account, you can use the resource's `buckets` attribute alongside `.all()`, which gives you the complete list of `Bucket` instances:

Python

>>>

```
>>> for bucket in s3_resource.buckets.all():
...     print(bucket.name)
...
firstpythonbucket7250e773-c4b1-422a-b51f-c45a52af9304
secondpythonbucket2d5d99c5-ab96-4c30-b7f7-443a95f72644
```

You can use the `client` to retrieve the bucket information as well, but the code is more complex, as you need to extract it from the dictionary that the `client` returns:

Python

>>>

```
>>> for bucket_dict in s3_resource.meta.client.list_buckets().get('Buckets'):
...     print(bucket_dict['Name'])
...
firstpythonbucket7250e773-c4b1-422a-b51f-c45a52af9304
secondpythonbucket2d5d99c5-ab96-4c30-b7f7-443a95f72644
```

You have seen how to iterate through the buckets you have in your account. In the upcoming section, you'll pick one of your buckets and iteratively view the objects it contains.

Object Traversal

If you want to list all the objects from a bucket, the following code will generate an iterator for you:

Python

>>>

```
>>> for obj in first_bucket.objects.all():
...     print(obj.key)
...
127367firstfile.txt
616abesecoundfile.txt
fb937cthirdfile.txt
```

The `obj` variable is an `ObjectSummary`. This is a lightweight representation of an `Object`. The summary version doesn't support all of the attributes that the `Object` has. If you need to access them, use the `object()` sub-resource to create a new reference to the underlying stored key. Then you'll be able to extract the missing attributes:

Python

>>>

```
>>> for obj in first_bucket.objects.all():
...     subsrcc = obj.Object()
...     print(obj.key, obj.storage_class, obj.last_modified,
...           subsrcc.version_id, subsrcc.metadata)
...
127367firstfile.txt STANDARD 2018-10-05 15:09:46+00:00 eQgH6IC1VGcn7eXZ_.ayqm6Ndjjh0ADv {}
616abesecoundfile.txt STANDARD 2018-10-05 15:09:47+00:00 WIAExRLmoksJzLhN7jU5YzoJxYSu6Ey6 {}
fb937cthirdfile.txt STANDARD_IA 2018-10-05 15:09:05+00:00 null {}
```

You can now iteratively perform operations on your buckets and objects. You're almost done. There's one more thing you should know at this stage: how to delete all the resources you've created in this tutorial.

Deleting Buckets and Objects

To remove all the buckets and objects you have created, you must first make sure that your buckets have no objects

within them.

Deleting a Non-empty Bucket

To be able to delete a bucket, you must first delete every single object within the bucket, or else the `BucketNotEmpty` exception will be raised. When you have a versioned bucket, you need to delete every object and all its versions.

If you find that a LifeCycle rule that will do this automatically for you isn't suitable to your needs, here's how you can programatically delete the objects:

Python

```
def delete_all_objects(bucket_name):
    res = []
    bucket=s3_resource.Bucket(bucket_name)
    for obj_version in bucket.object_versions.all():
        res.append({'Key': obj_version.object_key,
                    'VersionId': obj_version.id})
    print(res)
    bucket.delete_objects(Delete={'Objects': res})
```

The above code works whether or not you have enabled versioning on your bucket. If you haven't, the version of the objects will be null. You can batch up to 1000 deletions in one API call, using `.delete_objects()` on your `Bucket` instance, which is more cost-effective than individually deleting each object.

Run the new function against the first bucket to remove all the versioned objects:

Python

>>>

```
>>> delete_all_objects(first_bucket_name)
[{'Key': '127367firstfile.txt', 'VersionId': 'eQgH6IC1VGcn7eXZ_.ayqm6Ndjjh0ADv'}, {'Key': '127367firstfile.tx
```

As a final test, you can upload a file to the second bucket. This bucket doesn't have versioning enabled, and thus the version will be null. Apply the same function to remove the contents:

Python

>>>

```
>>> s3_resource.Object(second_bucket_name, first_file_name).upload_file(
...     first_file_name)
>>> delete_all_objects(second_bucket_name)
[{'Key': '9c8b44firstfile.txt', 'VersionId': 'null'}]
```

You've successfully removed all the objects from both your buckets. You're now ready to delete the buckets.

Deleting Buckets

To finish off, you'll use `.delete()` on your `Bucket` instance to remove the first bucket:

Python

```
s3_resource.Bucket(first_bucket_name).delete()
```

If you want, you can use the `client` version to remove the second bucket:

Python

```
s3_resource.meta.client.delete_bucket(Bucket=second_bucket_name)
```

Both the operations were successful because you emptied each bucket before attempting to delete it.

You've now run some of the most important operations that you can perform with S3 and Boto3. Congratulations on

making it this far! As a bonus, let's explore some of the advantages of managing S3 resources with Infrastructure as Code.

Python Code or Infrastructure as Code (IaC)?

As you've seen, most of the interactions you've had with S3 in this tutorial had to do with objects. You didn't see many bucket-related operations, such as adding policies to the bucket, adding a LifeCycle rule to transition your objects through the storage classes, archive them to Glacier or delete them altogether or enforcing that all objects be encrypted by configuring Bucket Encryption.

Manually managing the state of your buckets via Boto3's clients or resources becomes increasingly difficult as your application starts adding other services and grows more complex. To monitor your infrastructure in concert with Boto3, consider using an Infrastructure as Code (IaC) tool such as CloudFormation or Terraform to manage your application's infrastructure. Either one of these tools will maintain the state of your infrastructure and inform you of the changes that you've applied.

If you decide to go down this route, keep the following in mind:

- Any bucket related-operation that modifies the bucket in any way should be done via IaC.
- If you want all your objects to act in the same way (all encrypted, or all public, for example), usually there is a way to do this directly using IaC, by adding a Bucket Policy or a specific Bucket property.
- Bucket read operations, such as iterating through the contents of a bucket, should be done using Boto3.
- Object-related operations at an individual object level should be done using Boto3.

Conclusion

Congratulations on making it to the end of this tutorial!

You're now equipped to start working programmatically with S3. You now know how to create objects, upload them to S3, download their contents and change their attributes directly from your script, all while avoiding common pitfalls with Boto3.

May this tutorial be a stepping stone in your journey to building something great using AWS!

Further Reading

If you want to learn more, check out the following:

- [Boto3 documentation](#)
- [Generating Random Data in Python \(Guide\)](#)
- [IAM Policies and Bucket Policies and ACLs](#)
- [Object Tagging](#)
- [LifeCycle Configurations](#)
- [Cross-Region Replication](#)
- [CloudFormation](#)
- [Terraform](#)

 **Watch Now** This tutorial has a related video course created by the Real Python team. Watch it together with the written tutorial to deepen your understanding: [Python, Boto3, and AWS S3: Demystified](#)



Get a short & sweet **Python Trick** delivered to your inbox every couple of days. No spam ever. Unsubscribe any time. Curated by the Real Python team.

```
1 # How to merge two dicts
2 # in Python 3.5+
3
4 >>> x = {'a': 1, 'b': 2}
5 >>> y = {'b': 3, 'c': 4}
6
7 >>> z = {**x, **y}
8
9 >>> z
10 {'c': 4, 'a': 1, 'b': 3}
```

Email Address

[Send Me Python Tricks »](#)

About Ralu Bolovan

Ralu is an avid Pythonista and writes for Real Python. She is a DevOps engineer specializing in cloud computing, with a penchant for AWS.

[» More about Ralu](#)

Each tutorial at Real Python is created by a team of developers so that it meets our high quality standards. The team members who worked on this tutorial are:

[Adriana](#)

[Brad](#)

[Joanna](#)

Master Real-World Python Skills With Unlimited Access to Real Python

— FREE Email Series —



```
1 # How to merge two dicts
2 # in Python 3.5+
3
4 >>> x = {'a': 1, 'b': 2}
5 >>> y = {'b': 3, 'c': 4}
6
7 >>> z = {**x, **y}
8
9 >>> z
10 {'c': 4, 'a': 1, 'b': 3}
```

What Do You

Email...

Get Python Tricks »

No spam. Unsubscribe any time.

Real Python Comment Policy: The most useful comments are those written with the goal of learning from or helping out other readers—after reading the whole article and all the earlier comments. Complaints and insults generally won't make the cut here.

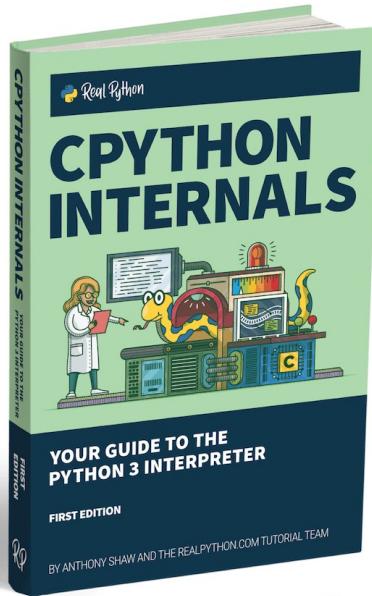
[All Tutorial Topics](#)

What's your #1 takeaway? [devops](#) [django](#) [docker](#) [flask](#) [front-end](#) [intermediate](#) [machine-learning](#) [found skills to use?](#)
Leave a comment below: [projects](#) [python](#) [testing](#) [tools](#) [web-dev](#) [web-scraping](#)

Keep Learning

Related Tutorial Categories: [devops](#) [intermediate](#)

Recommended Video Course: [Python, Boto3, and AWS S3: Demystified](#)



[Download a Free Chapter »](#)

Table of Contents

© 2012–2013

Py

- [Installation](#)
- [Client Versus Resource](#)
- [Common Operations](#)
- [Advanced Configurations](#)
- [Traversals](#)
- [Deleting Buckets and Objects](#)
- [Python Code or Infrastructure as Code \(IaC\)?](#)
- [Conclusion](#)
- [Further Reading](#)

stagram ·

ct

[Tweet](#) [Share](#) [Email](#)

Recommended Video Course

[Python, Boto3, and AWS S3: Demystified](#)



High Quality
Python Video Courses

[Watch Now »](#)